

```
In [1]: import pandas as pd
import numpy as np
import re
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from math import sqrt
from urllib.request import urlopen
import json
```

# Combined Cycle Power Plant Dataset

## Dataset Information:

(Extracted from UCI's Machine learning repository)

"The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011), when the power plant was set to work with full load. Features consist of hourly average ambient variables Temperature (T) [**°C**] , Ambient Pressure (AP) [**milibar**] , Relative Humidity (RH) [%] and Exhaust Vacuum (V) to predict the net hourly electrical energy output (EP) [**MW**] of the plant."

We will train a model to predict power output given three variables:

- Temperature (T) ..... [**°C**]
- Ambient Pressure (AP) ..... [**milibar**]
- Relative Humidity (RH) ..... [%]

Then we will take average data for these three variables from diferent locations in the US and feed them to the model, so we can predict in which locations a power plant would have the highest energy output.

```
In [2]: data = pd.read_excel('CCPP/Folds5x2_pp.xlsx')
data.head(10)
```

Out[2]:

	AT	V	AP	RH	PE
0	14.96	41.76	1024.07	73.17	463.26
1	25.18	62.96	1020.04	59.08	444.37
2	5.11	39.40	1012.16	92.14	488.56
3	20.86	57.32	1010.24	76.64	446.48
4	10.82	37.50	1009.23	96.62	473.90
5	26.27	59.44	1012.23	58.77	443.67
6	15.89	43.96	1014.02	75.24	467.35
7	9.48	44.71	1019.12	66.43	478.42
8	14.64	45.00	1021.78	41.25	475.98
9	11.74	43.56	1015.14	70.72	477.50

```
In [3]: data = data.drop(['V'], axis = 1)

data.columns = ['Temperature', 'Pressure', 'Humidity', 'Power Out']

data.head()
```

Out[3]:

	Temperature	Pressure	Humidity	Power Out
0	14.96	1024.07	73.17	463.26
1	25.18	1020.04	59.08	444.37
2	5.11	1012.16	92.14	488.56
3	20.86	1010.24	76.64	446.48
4	10.82	1009.23	96.62	473.90

```
In [4]: data.describe()
```

Out[4]:

	Temperature	Pressure	Humidity	Power Out
<b>count</b>	9568.000000	9568.000000	9568.000000	9568.000000
<b>mean</b>	19.651231	1013.259078	73.308978	454.365009
<b>std</b>	7.452473	5.938784	14.600269	17.066995
<b>min</b>	1.810000	992.890000	25.560000	420.260000
<b>25%</b>	13.510000	1009.100000	63.327500	439.750000
<b>50%</b>	20.345000	1012.940000	74.975000	451.550000
<b>75%</b>	25.720000	1017.260000	84.830000	468.430000
<b>max</b>	37.110000	1033.300000	100.160000	495.760000

```
In [5]: data.corr()
```

Out[5]:

	Temperature	Pressure	Humidity	Power Out
<b>Temperature</b>	1.000000	-0.507549	-0.542535	-0.948128
<b>Pressure</b>	-0.507549	1.000000	0.099574	0.518429
<b>Humidity</b>	-0.542535	0.099574	1.000000	0.389794
<b>Power Out</b>	-0.948128	0.518429	0.389794	1.000000

## Weather Datasets

We are going to extract information about the average Temperature, Relative Humidity and Ambient Pressure in US from three different datasets with yearly data taken from different stations around US:

- US Weather Stations Temperatures Dataset.
- US Weather Stations Relative Humidity Dataset.
- Station metadata Dataset.

The last one contains information about stations all over the world. This includes elevation over sea level. With this we can estimate the Ambient Pressure. We'll extract only the data for the stations in the US.

## Temperature Data

```
In [6]: tempdata = pd.read_excel('Weather/temp.xlsx')
tempdata.head()
```

Out[6]:

	wban	city	st	YRS	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
0	13876	BIRMINGHAM	AL	30	43.8	47.7	55.2	62.5	70.6	77.7	81.1	80.7	74.7	64.7
1	3856	HUNTSVILLE	AL	30	41.5	45.7	53.5	61.8	70.3	77.7	80.6	80.1	73.7	62.1
2	13894	MOBILE	AL	30	50.4	53.8	60.2	66.4	74.1	79.8	81.8	81.6	77.5	68.4
3	13895	MONTGOMERY	AL	30	46.6	50.5	57.5	64.1	72.4	79.0	81.8	81.5	76.3	65.1
4	26451	ANCHORAGE	AK	30	17.1	20.2	26.6	36.8	47.8	55.2	58.8	56.7	48.6	34.1

```
In [7]: tempdata = tempdata[['wban', 'city', 'st', 'ANN']]
tempdata.columns = ['Wban', 'City', 'State', 'Temp(°F)']
tempdata.head()
```

Out[7]:

	Wban	City	State	Temp(°F)
0	13876	BIRMINGHAM	AL	63.3
1	3856	HUNTSVILLE	AL	62.1
2	13894	MOBILE	AL	67.2
3	13895	MONTGOMERY	AL	65.1
4	26451	ANCHORAGE	AK	37.1

## Fahrenheit to Celsius conversion

$$T(^{\circ}\text{C}) = (T(^{\circ}\text{F}) - 32) / 1.8$$

```
In [8]: tempF = np.array(tempdata['Temp(°F)'])
temperature = (tempF - 32)/ 1.8
tempdata['Temperature'] = temperature
tempdata = tempdata.drop('Temp(°F)', axis = 1)
tempdata.head()
```

Out[8]:

	Wban	City	State	Temperature
0	13876	BIRMINGHAM	AL	17.388889
1	3856	HUNTSVILLE	AL	16.722222
2	13894	MOBILE	AL	19.555556
3	13895	MONTGOMERY	AL	18.388889
4	26451	ANCHORAGE	AK	2.833333

```
In [9]: tempdata.isnull().any()
```

```
Out[9]: Wban          False
City          False
State         False
Temperature   False
dtype: bool
```

```
In [10]: tempdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 263 entries, 0 to 262
Data columns (total 4 columns):
Wban          263 non-null int64
City          263 non-null object
State         263 non-null object
Temperature   263 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 8.3+ KB
```

## Humidity Data

```
In [11]: humdata = pd.read_excel('Weather/relhum.xlsx')
humdata.head()
```

Out[11]:

		Unnamed: 0	POR	JAN	Unnamed: 3	FEB	Unnamed: 5	MAR	Unnamed: 7	APR
0		NaN	NaN	M	A	M	A	M	A	M
1	13876	BIRMINGHAM,AL	194801-201812	81	60	80	56	80	52	83
2	03856	HUNTSVILLE,AL	195809-201812	81	64	81	60	80	55	82
3	13894	MOBILE,AL	194801-201812	83	60	84	58	85	54	88
4	13895	MONTGOMERY,AL	194801-201812	83	59	82	56	83	52	87

5 rows x 28 columns

```
In [12]: humdata.columns = [ '0', 'POR', 'JAN', 'JAN', 'FEB', 'FEB', 'MAR',
                              'MAR', 'APR', 'APR', 'MAY', 'MAY', 'JUN',
                              'JUN', 'JUL', 'JUL', 'AUG', 'AUG', 'SEP',
                              'SEP', 'OCT', 'OCT', 'NOV', 'NOV', 'DEC',
                              'DEC', 'ANN1', 'ANN2' ]
```

```
In [13]: humdata.head()
```

Out[13]:

		0	POR	JAN	JAN	FEB	FEB	MAR	MAR	APR	APR	...	SEP	!
0		NaN	NaN	M	A	M	A	M	A	M	A	...	M	
1	13876	BIRMINGHAM,AL	194801-201812	81	60	80	56	80	52	83	50	...	87	
2	03856	HUNTSVILLE,AL	195809-201812	81	64	81	60	80	55	82	51	...	89	
3	13894	MOBILE,AL	194801-201812	83	60	84	58	85	54	88	53	...	90	
4	13895	MONTGOMERY,AL	194801-201812	83	59	82	56	83	52	87	51	...	90	

5 rows x 28 columns

```
In [14]: humdata = humdata[['0' , 'ANN1' , 'ANN2']].drop([0])
humdata.head()
```

Out[14]:

	0	ANN1	ANN2
1	13876BIRMINGHAM,AL	84	55
2	03856HUNTSVILLE,AL	85	57
3	13894MOBILE,AL	87	57
4	13895MONTGOMERY,AL	87	55
5	26451ANCHORAGE,AK	74	64

```
In [15]: wban = humdata['0']
wban = [re.findall(r'[0-9]+', i)[0] for i in wban]
```

```
In [16]: humdata['Wban'] = wban
humdata = humdata.drop(['0'], axis = 1)

humdata['Humidity'] = humdata[['ANN1' , 'ANN2']].mean(axis=1)
humdata = humdata.drop(['ANN1' , 'ANN2'], axis = 1)
humdata.head()
```

Out[16]:

	Wban	Humidity
1	13876	69.5
2	03856	71.0
3	13894	72.0
4	13895	71.0
5	26451	69.0

```
In [17]: humdata['Wban'] = humdata['Wban'].astype('int64')
```

```
In [18]: humdata.isnull().any()
```

```
Out[18]: Wban          False
Humidity        False
dtype: bool
```

```
In [19]: humdata.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 266 entries, 1 to 266
Data columns (total 2 columns):
Wban          266 non-null int64
Humidity      266 non-null float64
dtypes: float64(1), int64(1)
memory usage: 6.2 KB
```

```
In [20]: print (humdata['Wban'].dtype)
print (tempdata['Wban'].dtype)

int64
int64
```

## Pressure Data

```
In [21]: presdata = pd.read_excel('Weather/stations.xlsx')
presdata.head()
```

Out[21]:

	Unnamed: 0	Unnamed: 1	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnamed: 5	Unnamed: 6	Unnamed: 7
0	USAF	WBAN	STATION NAME	CTRY	ST CALL	LAT	LON	ELEV(M)
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	7018	99999	WXPOD 7018	NaN	NaN	0	0	7018
3	7026	99999	WXPOD 7026	AF	NaN	0	0	7026
4	7070	99999	WXPOD 7070	AF	NaN	0	0	7070



```
In [22]: columns = presdata.iloc[0]
presdata.columns = columns

presdata = presdata.drop([0])
presdata.head()
```

Out[22]:

	USAF	WBAN	STATION NAME	CTRY	ST CALL	LAT	LON	ELEV(M)	BEGIN	END
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	7018	99999	WXPOD 7018	NaN	NaN	0	0	7018	20110309	20130730
3	7026	99999	WXPOD 7026	AF	NaN	0	0	7026	20120713	20170822
4	7070	99999	WXPOD 7070	AF	NaN	0	0	7070	20140923	20150926
5	8260	99999	WXPOD8270	NaN	NaN	0	0	0	20050101	20100920

```
In [23]: presdata = presdata[presdata['CTRY']=='US']
presdata.head()
```

Out[23]:

	USAF	WBAN	STATION NAME	CTRY	ST CALL	LAT	LON	ELEV(M)	BEGIN	END
13135	621010	99999	MOORED BUOY	US	NaN	50.6	-2.933	-999	20080721	20080721
13137	621110	99999	MOORED BUOY	US	NaN	58.9	-0.2	-999	20041118	20041118
13138	621130	99999	MOORED BUOY	US	NaN	58.4	0.3	-999	20040726	20040726
13139	621160	99999	MOORED BUOY	US	NaN	58.1	1.8	-999	20040829	20040829
13140	621170	99999	MOORED BUOY	US	NaN	57.9	0.1	-999	20040726	20040726

```
In [24]: presdata = presdata[['WBAN', 'ELEV(M)']]

presdata['ELEV(M)'] = pd.to_numeric(presdata['ELEV(M)'], errors='coerce')

presdata['WBAN'] = presdata['WBAN'].astype('int64')

presdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7326 entries, 13135 to 29745
Data columns (total 2 columns):
WBAN          7326 non-null int64
ELEV(M)       7209 non-null float64
dtypes: float64(1), int64(1)
memory usage: 171.7 KB
```

```
In [25]: presdata.dropna(inplace=True)

presdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7209 entries, 13135 to 29745
Data columns (total 2 columns):
WBAN          7209 non-null int64
ELEV(M)       7209 non-null float64
dtypes: float64(1), int64(1)
memory usage: 169.0 KB
```

```
In [26]: presdata.columns = ['Wban', 'Elevation']
presdata.head()
```

Out[26]:

	Wban	Elevation
13135	99999	-999.0
13137	99999	-999.0
13138	99999	-999.0
13139	99999	-999.0
13140	99999	-999.0

```
In [27]: presdata = presdata[presdata['Elevation']>=0]
presdata.head()
```

Out[27]:

	Wban	Elevation
13149	99999	0.0
14465	99999	1224.0
14466	93218	317.0
14467	99999	317.0
14468	93217	43.0

## Average ambient pressure calculation

For this calculation we will use the following expresion:

$$p \sim 101325 (1 - 2.25577 \cdot 10^{-5} \cdot h)^5$$

Where:

101325 = normal temperature and pressure at sea level (Pa)

p = air pressure (Pa)

h = altitude above sea level (m)

Pascal to milibar conversion:

$$1 \text{ Pa} = 0.01 \text{ milibar}$$

```
In [28]: elev = np.array(presdata['Elevation'])

Pasc = 101325*(1 - 2.25577*10**(-5) * elev)**5

milibar = Pasc * 0.01
```

```
In [29]: presdata['Pressure'] = milibar
presdata = presdata.drop('Elevation', axis = 1)

presdata.head()
```

Out[29]:

	Wban	Pressure
<b>13149</b>	99999	1013.250000
<b>14465</b>	99999	880.881800
<b>14466</b>	93218	977.536727
<b>14467</b>	99999	977.536727
<b>14468</b>	93217	1008.345357

```
In [30]: presdata['Wban'].value_counts().head()
```

Out[30]:

99999	3548
23176	6
3849	5
14897	4
24255	4

Name: Wban, dtype: int64

```
In [31]: presdata[presdata['Wban']==23176]
```

Out[31]:

	Wban	Pressure
<b>19714</b>	23176	849.441660
<b>19722</b>	23176	849.650098
<b>19743</b>	23176	849.650098
<b>21748</b>	23176	849.650098
<b>21749</b>	23176	849.650098
<b>29011</b>	23176	849.650098

```
In [32]: presdata['Wban'] = presdata[presdata['Wban']!=99999]
presdata = presdata.drop_duplicates(subset = 'Wban')
presdata['Wban'].value_counts().head()
```

Out[32]:

54791.0	1
24164.0	1
24048.0	1
4898.0	1
3930.0	1

Name: Wban, dtype: int64

# Merging weather datasets into US Locations

```
In [33]: USloc = tempdata.merge(humdata, how='inner')
print (len(USloc))
USloc.head(10)
```

260

Out[33]:

	Wban	City	State	Temperature	Humidity
0	13876	BIRMINGHAM	AL	17.388889	69.5
1	3856	HUNTSVILLE	AL	16.722222	71.0
2	13894	MOBILE	AL	19.555556	72.0
3	13895	MONTGOMERY	AL	18.388889	71.0
4	26451	ANCHORAGE	AK	2.833333	69.0
5	25308	ANNETTE	AK	8.111111	77.5
6	27502	BARROW	AK	-11.222222	81.0
7	26615	BETHEL	AK	-0.722222	78.5
8	26533	BETTLES	AK	-4.722222	67.0
9	26415	BIG DELTA	AK	-1.666667	62.5

```
In [34]: usloc = USloc.merge(presdata, how='left')

print (len(usloc))
usloc.head(10)
```

260

Out[34]:

	Wban	City	State	Temperature	Humidity	Pressure
0	13876	BIRMINGHAM	AL	17.388889	69.5	992.002445
1	3856	HUNTSVILLE	AL	16.722222	71.0	991.699106
2	13894	MOBILE	AL	19.555556	72.0	1005.786554
3	13895	MONTGOMERY	AL	18.388889	71.0	1006.229708
4	26451	ANCHORAGE	AK	2.833333	69.0	1009.074145
5	25308	ANNETTE	AK	8.111111	77.5	1009.461485
6	27502	BARROW	AK	-11.222222	81.0	1012.164777
7	26615	BETHEL	AK	-0.722222	78.5	1009.700784
8	26533	BETTLES	AK	-4.722222	67.0	991.362149
9	26415	BIG DELTA	AK	-1.666667	62.5	969.545252

```
In [35]: usloc.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 260 entries, 0 to 259
Data columns (total 6 columns):
Wban          260 non-null int64
City          260 non-null object
State         260 non-null object
Temperature   260 non-null float64
Humidity      260 non-null float64
Pressure      251 non-null float64
dtypes: float64(3), int64(1), object(2)
memory usage: 14.2+ KB
```

```
In [36]: usloc = usloc.dropna()

usloc.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 251 entries, 0 to 250
Data columns (total 6 columns):
Wban          251 non-null int64
City          251 non-null object
State         251 non-null object
Temperature   251 non-null float64
Humidity      251 non-null float64
Pressure      251 non-null float64
dtypes: float64(3), int64(1), object(2)
memory usage: 13.7+ KB
```

## Multiple Linear Regression Model

We are going to create a model to predict energy output (EP) [MW], taking Temperature (T) [°C] , Ambient Pressure (AP) [milibar] , Relative Humidity (RH) [%] as parameters.

We will use the dataset called 'data', which contains the values for the CCGP.

We will use scikit-learn, and train the model as a multiple linear regression

```
In [37]: X = data[['Temperature', 'Pressure', 'Humidity']]
Y = data[['Power Out']]
```

```
In [38]: Xtrain, Xtest, Ytrain, Ytest = train_test_split(X,Y, test_size=0.25,
random_state=10)
```

```
In [39]: regression = LinearRegression()
model = regression.fit(Xtrain, Ytrain)
```

```
In [40]: Ypred = model.predict(Xtest)
Ypred
```

```
Out[40]: array([[455.39353049],
[447.47658286],
[437.21191769],
...,
[487.93736352],
[440.41536484],
[453.40714577]])
```

```
In [41]: RMSE = sqrt(mean_squared_error(y_true = Ytest, y_pred = Ypred))
RMSE
```

```
Out[41]: 4.657571391074255
```

```
In [42]: r2 = r2_score(Ytest, Ypred)
r2
```

```
Out[42]: 0.9266765701220129
```

## Performance prediction at US locations

In this step we will use the model previously trained with the CCPP data, to predict the expected power output of a CCPP situated in different US locations based on their climate average conditions.

```
In [43]: usweather = usloc[['Temperature', 'Pressure', 'Humidity']]

usweather.head()
```

```
Out[43]:
```

	Temperature	Pressure	Humidity
0	17.388889	992.002445	69.5
1	16.722222	991.699106	71.0
2	19.555556	1005.786554	72.0
3	18.388889	1006.229708	71.0
4	2.833333	1009.074145	69.0

```
In [44]: X.head()
```

```
Out[44]:
```

	Temperature	Pressure	Humidity
0	14.96	1024.07	73.17
1	25.18	1020.04	59.08
2	5.11	1012.16	92.14
3	20.86	1010.24	76.64
4	10.82	1009.23	96.62



```
In [45]: USpred = model.predict(usweather)
USpred[:5]
```

```
Out[45]: array([[459.91897085],
                [461.1913595 ],
                [454.64134399],
                [457.62249043],
                [495.01181029]])
```

```
In [46]: usloc['Power_out'] = USpred

usloc.head()
```

```
Out[46]:
```

	Wban	City	State	Temperature	Humidity	Pressure	Power_out
0	13876	BIRMINGHAM	AL	17.388889	69.5	992.002445	459.918971
1	3856	HUNTSVILLE	AL	16.722222	71.0	991.699106	461.191360
2	13894	MOBILE	AL	19.555556	72.0	1005.786554	454.641344
3	13895	MONTGOMERY	AL	18.388889	71.0	1006.229708	457.622490
4	26451	ANCHORAGE	AK	2.833333	69.0	1009.074145	495.011810

## Findings

```
In [47]: usloc.Power_out.describe()
```

```
Out[47]: count    251.000000
mean      472.345828
std       14.272938
min       440.563445
25%      462.034303
50%      473.768200
75%      480.319051
max       526.034436
Name: Power_out, dtype: float64
```

```
In [48]: usloc.drop('Wban',axis=1).corr()
```

```
Out[48]:
```

	Temperature	Humidity	Pressure	Power_out
Temperature	1.000000	-0.057025	0.236537	-0.991844
Humidity	-0.057025	1.000000	0.586223	-0.048887
Pressure	0.236537	0.586223	1.000000	-0.243565
Power_out	-0.991844	-0.048887	-0.243565	1.000000

```
In [49]: usloc.sort_values('Power_out', ascending=False).head(10)
```

```
Out[49]:
```

	Wban	City	State	Temperature	Humidity	Pressure	Power_out
6	27502	BARROW	AK	-11.222222	81.0	1012.164777	526.034436
147	14755	MT. WASHINGTON	NH	-2.555556	0.0	812.822902	516.454667
8	26533	BETTLES	AK	-4.722222	67.0	991.362149	512.871904
17	26616	KOTZEBUE	AK	-5.055556	78.0	1012.210452	512.005739
11	26411	FAIRBANKS	AK	-2.388889	65.0	998.288099	507.921088
18	26510	MCGRATH	AK	-2.555556	68.5	1001.703277	507.703816
19	26617	NOME	AK	-2.555556	76.0	1012.792951	506.490944
12	26425	GULKANA	AK	-2.111111	66.0	959.996103	506.043838
9	26415	BIG DELTA	AK	-1.666667	62.5	969.545252	505.946333
7	26615	BETHEL	AK	-0.722222	78.5	1009.700784	501.556013

```
In [50]: usloc.sort_values('Power_out', ascending=False).tail(10)
```

```
Out[50]:
```

	Wban	City	State	Temperature	Humidity	Pressure	Power_out
61	12842	TAMPA	FL	23.000000	72.0	1012.587332	446.648664
70	21504	HILO	HI	23.277778	74.0	1011.925011	445.569949
211	12919	BROWNSVILLE	TX	23.611111	75.0	1012.416009	444.591033
53	12835	FORT MYERS	FL	23.944444	71.5	1012.724408	444.511729
72	22516	KAHULUI	HI	24.388889	67.5	1011.479853	444.228015
63	12844	WEST PALM BEACH	FL	24.111111	71.5	1012.587332	444.112604
73	22536	LIHUE	HI	24.333333	72.0	1009.769163	443.409941
71	22521	HONOLULU	HI	25.388889	65.0	1013.010029	442.398239
57	12839	MIAMI	FL	25.111111	71.5	1012.244709	441.730603
56	12836	KEY WEST	FL	25.444444	73.5	1013.215716	440.563445

```
In [51]: statemean = usloc.groupby('State', as_index=False).mean()
statemean = statemean.sort_values('Power_out', ascending=False)
statemean.head()
```

Out[51]:

	State	Wban	Temperature	Humidity	Pressure	Power_out
30	NH	14750.00	2.722222	33.250	907.087592	499.751931
0	AK	26143.45	0.519444	74.325	1003.855505	499.293539
28	ND	36963.75	5.416667	70.875	968.073548	487.415893
21	ME	14685.50	6.222222	69.250	1001.692198	486.723986
23	MN	14920.80	5.755556	71.900	973.431288	486.548040

## Plotting color US map

```
In [52]: city = pd.read_excel('uscities.xlsx')
city.head(10)
```

Out[52]:

	city	city_ascii	state_id	state_name	county_fips	county_name	county_fips_all	co
0	South Creek	South Creek	WA	Washington	53053	Pierce	53053	
1	Roslyn	Roslyn	WA	Washington	53037	Kittitas	53037	
2	Sprague	Sprague	WA	Washington	53043	Lincoln	53043	
3	Gig Harbor	Gig Harbor	WA	Washington	53053	Pierce	53053	
4	Lake Cassidy	Lake Cassidy	WA	Washington	53061	Snohomish	53061	
5	Tenino	Tenino	WA	Washington	53067	Thurston	53067	
6	Jamestown	Jamestown	WA	Washington	53009	Clallam	53009	
7	Three Lakes	Three Lakes	WA	Washington	53061	Snohomish	53061	
8	Curlew Lake	Curlew Lake	WA	Washington	53019	Ferry	53019	
9	Chain Lake	Chain Lake	WA	Washington	53061	Snohomish	53061	

```
In [53]: city = city[['city', 'state_id', 'county_fips']]
city.head()
```

Out[53]:

	city	state_id	county_fips
0	South Creek	WA	53053
1	Roslyn	WA	53037
2	Sprague	WA	53043
3	Gig Harbor	WA	53053
4	Lake Cassidy	WA	53061

```
In [54]: city.city = city.city.str.upper()
city.head()
```

Out[54]:

	city	state_id	county_fips
0	SOUTH CREEK	WA	53053
1	ROSLYN	WA	53037
2	SPRAGUE	WA	53043
3	GIG HARBOR	WA	53053
4	LAKE CASSIDY	WA	53061

```
In [55]: DF = usloc[['City', 'State', 'Power_out']]
DF.columns = ['city', 'state_id', 'Power_out']
DF.city = DF.city.str.rstrip()
DF.head()
```

/Users/sandra/anaconda3/lib/python3.7/site-packages/pandas/core/generic.py:5096: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>  
self[name] = value

Out[55]:

	city	state_id	Power_out
0	BIRMINGHAM	AL	459.918971
1	HUNTSVILLE	AL	461.191360
2	MOBILE	AL	454.641344
3	MONTGOMERY	AL	457.622490
4	ANCHORAGE	AK	495.011810

```
In [56]: city.head()
```

Out[56]:

	city	state_id	county_fips
0	SOUTH CREEK	WA	53053
1	ROSLYN	WA	53037
2	SPRAGUE	WA	53043
3	GIG HARBOR	WA	53053
4	LAKE CASSIDY	WA	53061

```
In [57]: city.city[0]
```

Out[57]: 'SOUTH CREEK'

n=0 for i in city.city: if 'Huntsville' in i: print (city.iloc[n]) n =n+1

```
In [58]: x = pd.merge(Df, city, on=['city', 'state_id'])
x.head()
```

Out[58]:

	city	state_id	Power_out	county_fips
0	BIRMINGHAM	AL	459.918971	1073
1	HUNTSVILLE	AL	461.191360	1089
2	MOBILE	AL	454.641344	1097
3	MONTGOMERY	AL	457.622490	1101
4	ANCHORAGE	AK	495.011810	2020

```
In [59]: x.county_fips[x.county_fips<10000]= ('0'+ x.county_fips.astype(str))
```

```
/Users/sandra/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
"""Entry point for launching an IPython kernel.
```

```
In [60]: len (x)
```

Out[60]: 232

```
In [61]: x.head()
```


Out[61]:

	city	state_id	Power_out	county_fips
0	BIRMINGHAM	AL	459.918971	01073
1	HUNTSVILLE	AL	461.191360	01089
2	MOBILE	AL	454.641344	01097
3	MONTGOMERY	AL	457.622490	01101
4	ANCHORAGE	AK	495.011810	02020

```
In [62]: with urlopen('https://raw.githubusercontent.com/plotly/datasets/master/geojson-counties-fips.json') as response:
    counties = json.load(response)

import plotly.express as px

fig = px.choropleth(x, geojson=counties, locations='county_fips', color=
r='Power_out',
                    color_continuous_scale="Viridis",
                    range_color=(439, 512),
                    scope="usa", labels={'Power_out': 'Power Out
put [MW]'},
                    )
fig.update_layout(margin={"r":0,"t":0,"l":0,"b":0})
fig.show()
```



# **Combined Cycle Power Plant Performance**

Sandra Poza



# Abstract

This project studies the relationship between weather conditions and performance of a combined cycle power plant (CCPP). We used a dataset containing CCPP performance to train a regression model that predict the power output of the plant.

Then with data from US weather datasets we predicted the performance it would have in different US locations.

It was found that it would perform better in coldest weather in northern areas.

# Motivation

Energy efficiency has become an essential matter in today's society. For climate, economic and sustainability reasons, every one gets benefit from a more effective use of energy.

In this project we try to find out where would a Combined Cycle Power Plant (CCPP) have the highest power output (which would mean a highest efficiency), based on the weather conditions.

# Dataset(s)

**Combined Cycle Power Plant Data Set:** The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011). Obtained from UCI's Machine learning repository.

**US Weather datasets:** Obtained from NOAA National Centers for Environmental Information

- **Temperatures dataset:** Contains monthly and yearly average temperatures for different US weather stations.
- **Relative Humidity dataset:** Contains monthly and yearly average relative humidity for different US weather stations.
- **Weather station locations dataset:** Contains weather stations information with elevation, used for ambient pressure calculation.

**US cities dataset:** Contains US cities information, including the fips codes. Obtained from <https://simplemaps.com>

# Methods

The Combined Cycle Power Plant Dataset contains data of a CCPP working at full load over 6 years. This include the following features:

Temperature (T) [°C]

Ambient Pressure (AP) [milibar]

Relative Humidity (RH) [%]

**Energy output (EP) [MW] of the plant**

With this data we will train a model that predict **Energy output** based on the 3 first. We will use scikit-learn, and train the model as a multiple linear regression.

Then, from the weather datasets we are going to extract information about the average Temperature, Relative Humidity and Ambient Pressure in different US locations.

Then we will use the model previously trained with the CCPP data, and feed the weather average data from different locations to predict the expected power output of a CCPP in this locations.

We will use Matplotlib and Plotly for plotting results.

# Findings

## Multiple Linear Regression Model Accuracy

The mean squared error for the model with test data is 4.65. Taking into consideration that the output of the model is in the order of 400-500 it is a good number.

```
RMSE = sqrt(mean_squared_error(y_true = Ytest, y_pred = Ypred))  
RMSE
```

```
4.657571391074255
```

A better stimator is the R2 score. This is a statistical measure of how well the regression predictions approximate the real data points. It ranges from 0-1 being 1 the best possible result.

Our R2 score is 0.92, what means that **the model is highly accurate** on test data.

```
from sklearn.metrics import r2_score
```

```
r2 = r2_score(Ytest, Ypred)  
r2
```

```
0.9266765701220129
```

# Findings

## Performance prediction at US locations

We can appreciate how the energy output is higher in the coldest areas like Alaska, and lower in the warmer ones like Florida.

	Wban	City	State	Temperature	Humidity	Pressure	Power_out
6	27502	BARROW	AK	-11.222222	81.0	1012.164777	526.034436
147	14755	MT. WASHINGTON	NH	-2.555556	0.0	812.822902	516.454667
8	26533	BETTLES	AK	-4.722222	67.0	991.362149	512.871904
17	26616	KOTZEBUE	AK	-5.055556	78.0	1012.210452	512.005739
11	26411	FAIRBANKS	AK	-2.388889	65.0	998.288099	507.921088
18	26510	MCGRATH	AK	-2.555556	68.5	1001.703277	507.703816
19	26617	NOME	AK	-2.555556	76.0	1012.792951	506.490944
12	26425	GULKANA	AK	-2.111111	66.0	959.996103	506.043838
9	26415	BIG DELTA	AK	-1.666667	62.5	969.545252	505.946333
7	26615	BETHEL	AK	-0.722222	78.5	1009.700784	501.556013


10 Highest Energy Output

	Wban	City	State	Temperature	Humidity	Pressure	Power_out
61	12842	TAMPA	FL	23.000000	72.0	1012.587332	446.648664
70	21504	HILO	HI	23.277778	74.0	1011.925011	445.569949
211	12919	BROWNSVILLE	TX	23.611111	75.0	1012.416009	444.591033
53	12835	FORT MYERS	FL	23.944444	71.5	1012.724408	444.511729
72	22516	KAHULUI	HI	24.388889	67.5	1011.479853	444.228015
63	12844	WEST PALM BEACH	FL	24.111111	71.5	1012.587332	444.112604
73	22536	LIHUE	HI	24.333333	72.0	1009.769163	443.409941
71	22521	HONOLULU	HI	25.388889	65.0	1013.010029	442.398239
57	12839	MIAMI	FL	25.111111	71.5	1012.244709	441.730603
56	12836	KEY WEST	FL	25.444444	73.5	1013.215716	440.563445

10 Lowest Energy Output

# Findings

## Performance prediction at US locations



This map shows the power output that it would be obtained by a CCPP situated in different US counties.

We can observe the relationship between southern, warmer, and more humid areas with lower power output

Also the northern, colder and drier areas with highest power output.

It should be mentioned that Alaska is represented on the bottom left, but is in fact the northernmost state, with the lowest temperatures, so we can appreciate how that region is the most yellowish.

# Limitations

We did not have enough weather data to study the whole US map.

The weather data is average, so it doesn't reflect seasonal changes.

The CCPP dataset measures was taken in a fixed location, so the pressure changes were not well represented.



# Conclusions

We can conclude, that the performance of a Combined Cycle Power Plant would be the best in the coldest and driest possible weathers.

A CCPP situated in Alaska would have the highest efficiency, and one located in Florida the lowest. That means that the same power plant would consume less gas in Alaska to produce the same power output.